

Portfolio 8: Metropolis-Hastings Algorithm

Kieran Morris

For this portfolio, we will implement the Metropolis-Hastings algorithm to sample from the posterior distribution of the parameters of a logistic regression model, our toy dataset for today is the Pima Indians Diabetes dataset from the `mlbench` package.

```
library(mlbench)
data("PimaIndiansDiabetes")
head(PimaIndiansDiabetes)
```

##	pregnant	glucose	pressure	triceps	insulin	mass	pedigree	age	diabetes
## 1	6	148	72	35	0	33.6	0.627	50	pos
## 2	1	85	66	29	0	26.6	0.351	31	neg
## 3	8	183	64	0	0	23.3	0.672	32	pos
## 4	1	89	66	23	94	28.1	0.167	21	neg
## 5	0	137	40	35	168	43.1	2.288	33	pos
## 6	5	116	74	0	0	25.6	0.201	30	neg

As you can see the dataset has 8 features and a binary response variable `diabetes`, which takes values `pos` and `neg`. Which we will probably change later on! We setup a logistic regression model:

Task 1 - Metropolis Hastings

$$Pr_{\alpha, \beta}(Y_i = 1) = \frac{1}{1 + e^{-\alpha - \beta^T x_i}}$$

with the following likelihood function:

$$L_n(\alpha, \beta) = \prod_{i=1}^n Pr_{\alpha, \beta}(Y_i = y_i)$$

and posterior distribution:

$$\pi(\alpha, \beta | y) \propto L_n(\alpha, \beta) \pi(\alpha, \beta)$$

the log posterior given by

$$\log \pi(\alpha, \beta | y) = \sum_{i=1}^n \left(y_i \log \frac{1}{1 + e^{-\alpha - \beta^T x_i}} + (1 - y_i) \log \frac{e^{-\alpha - \beta^T x_i}}{1 + e^{-\alpha - \beta^T x_i}} \right) + \pi(\alpha, \beta)$$

where $\pi(\alpha, \beta)$ is the prior.

1. Choice of Distribution

We will use the suggested choice of proposal distribution, a multivariate normal distribution at each z :

$$Q(z, dz') = \mathcal{N}_{p+1}(z, c\Sigma_n)$$

for tuning parameter $c > 0$ and

$$\mu_n = \arg \max_{(\alpha, \beta) \in \mathbb{R}^p} \log \pi(\alpha, \beta | y), \quad \Sigma_n = -(\mathbf{H}_n(\mu_q))^{-1}$$

where Σ_n is the hessian of the log posterior at μ_q and can be estimated by the covariance matrix of the coefficients under $\pi(\alpha, \beta | y_0)$.

To get the covariance matrix, we first fit a logistic regression model to the data and output it:

```
data <- PimaIndiansDiabetes
y <- as.numeric((data$diabetes))-1
X <- scale(data[,1:8])
fit <- glm(y ~ X, family = "binomial")
z_0 <- fit$coefficients
Sigma_n <- summary(fit)$cov.scaled
```

2. Implementing Metropolis Hastings

We attempt to manually perform the Metropolis-Hastings algorithm. So we begin by coding up a function which produces the posterior likelihood, with a flat prior:

```
library(mvtnorm)
posterior_likelihood <- function(par, X, y){
  alpha <- par[1]
  beta <- par[2:9]
  p <- 1 - 1 / (1 + exp(alpha + beta%*t(X)))
  likelihood <- exp(sum(dbinom(y, size=1, prob=p, log=TRUE)) )
  return(likelihood)
}
```

Since we now have a proposal distribution and a posterior likelihood, we can now run the MH algorithm:

```
metropolis_hastings <- function(z_0, c, sigma_n, tmax, X, y){
  zs <- matrix(NA, nrow = tmax, ncol = length(z_0))
  z_current <- z_0
  z_current_ll <- posterior_likelihood(z_current, X, y)
  sigma <- c * sigma_n
  accepted <- 0
  for (i in 1:tmax){
    z_new <- rmvnorm(1, z_current, sigma)
    z_new_ll <- posterior_likelihood(z_new, X, y)
    alpha <- log(z_new_ll) - log(z_current_ll)
    if (runif(1) < min(1, exp(alpha))) {
      z_current <- z_new
      z_current_ll <- z_new_ll
      accepted <- accepted + 1
    }
  }
}
```

```

    }
    zs[i,] <- z_current
  }
  return(list(zs = zs, acceptance_rate = accepted/tmax))
}

```

we begin with $c = 1$ and run the algorithm for 10000 iterations:

```

mh <- metropolis_hastings(z_0, 1, Sigma_n, 10000, X, y)
mh$acceptance_rate

```

```
## [1] 0.1794
```

obviously this isn't close to our ideal 0.234 acceptance rate, so we will need to try a range of c values to get a better acceptance rate:

```

c_values <- seq(0.1, 2, 0.1)
acceptance_rates <- numeric(length(c_values))
for (i in 1:length(c_values)){
  acceptance_rates[i] <- metropolis_hastings(z_0, c_values[i], Sigma_n, 10000, X, y)$acceptance_rate
}
#Output the c values and their acceptance error
cbind(c_values, acceptance_rates)

```

```

##      c_values acceptance_rates
## [1,]      0.1          0.6521
## [2,]      0.2          0.5250
## [3,]      0.3          0.4426
## [4,]      0.4          0.3780
## [5,]      0.5          0.3231
## [6,]      0.6          0.2754
## [7,]      0.7          0.2355
## [8,]      0.8          0.2088
## [9,]      0.9          0.1906
## [10,]     1.0          0.1757
## [11,]     1.1          0.1556
## [12,]     1.2          0.1366
## [13,]     1.3          0.1305
## [14,]     1.4          0.1112
## [15,]     1.5          0.1045
## [16,]     1.6          0.0986
## [17,]     1.7          0.0805
## [18,]     1.8          0.0795
## [19,]     1.9          0.0754
## [20,]     2.0          0.0627

```

Our best choice for c is then 0.8. We can then run the Metropolis-Hastings algorithm with this value of c :

```
mh <- metropolis_hastings(z_0, 0.8, Sigma_n, 10000, X, y)
```


3. Evaluating convergence with the optimal c

Let's check the convergence of our algorithm, we begin by looking at the trace plots and autocorrelation plots:

```
library(dplyr)
library(tidyr)
library(ggplot2)

colnames(mh$zs) <- c("intercept", colnames(data)[1:8])
zs_plot <- mh$zs %>%
  as_tibble() %>%
  mutate(iteration = 1:10000) %>%
  pivot_longer(cols = 1:9)

ggplot(zs_plot, aes(iteration, value)) +
  geom_line(colour = "#ff009d") +
  facet_wrap(vars(name))
```

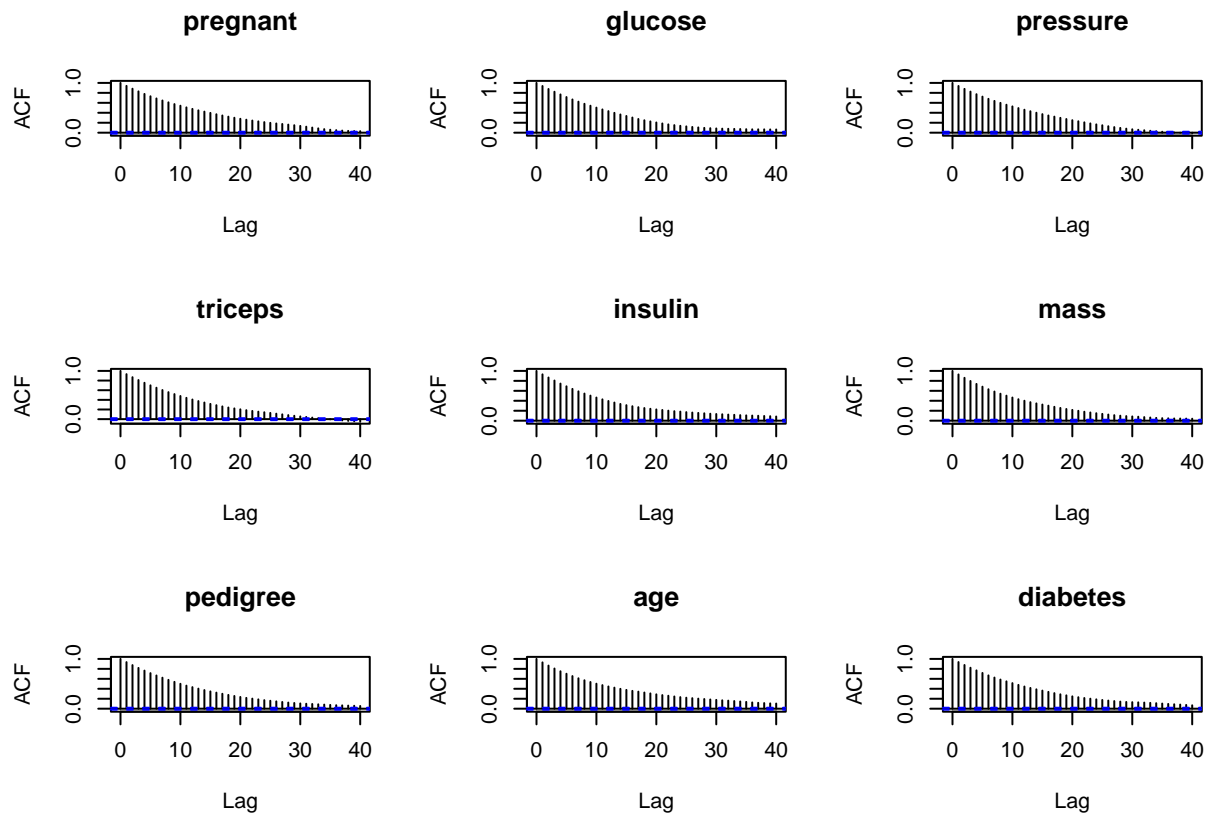


```
par(mfrow = c(3,3))
for (i in 1:9){
```

```

    acf(mh$zs[,i], main = colnames(data)[i])
  }

```



Since we did the work finding an optimal c value prior to modelling, we have a pretty strong acceptance rate and as we can see from the plots we have a tight distribution. The autocorrelation plots also are good indicators as they trail off pretty quickly.

5. Marginal posterior distributions

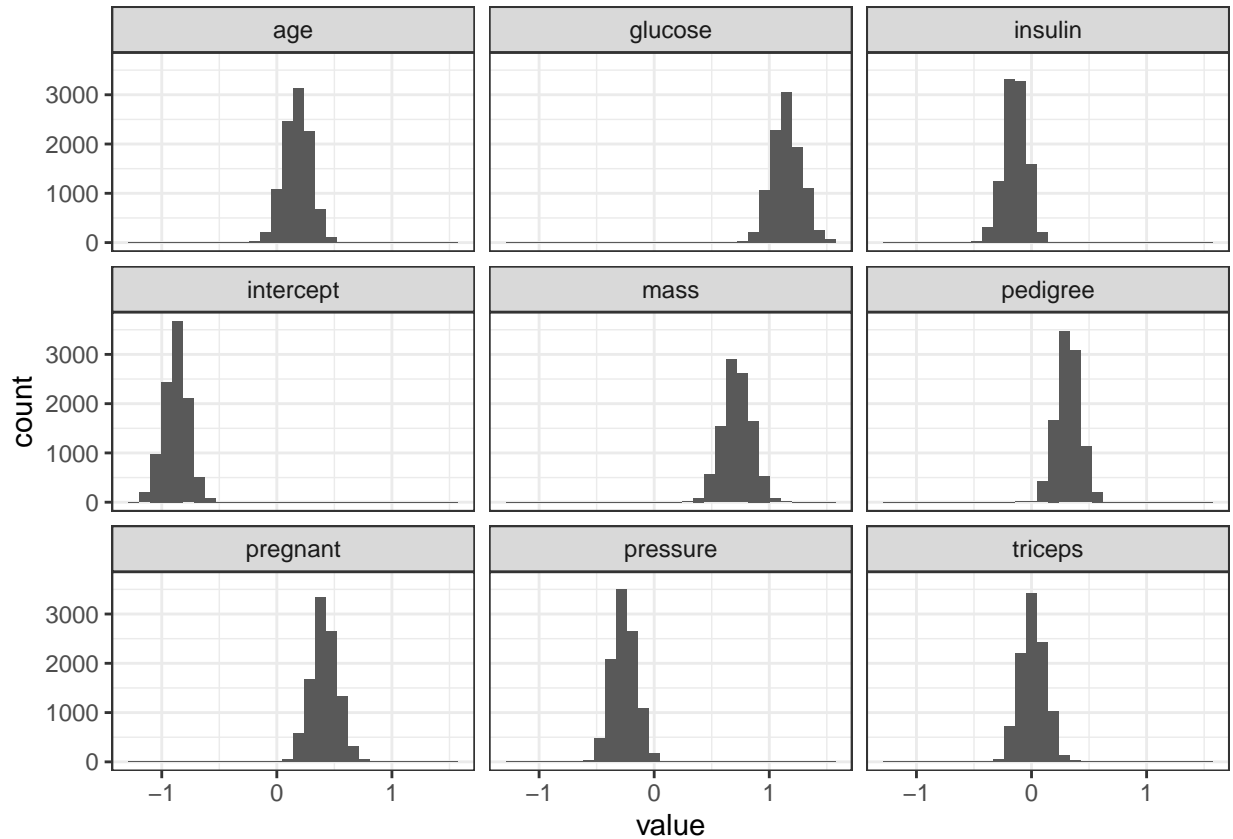
Now that we're happy with our model, we can plot the marginal posterior distributions for each parameter:

```

ggplot(zs_plot, aes(value)) +
  geom_histogram() +
  facet_wrap(vars(name))

```

'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.



Task 2 - Multiple Proposal Distributions

We now change our model to allow for multiple proposal distributions, we will use a mixture of normal distributions:

$$Q(z, dz') = \sum_{i=1}^k w_i \mathcal{N}_{p+1}(z, c_i \Sigma_n)$$

where w_i are the weights of the mixture and c_i are the tuning parameters. We will use $k = 3$ and $w_i = 1/3$ for all i .

1. Implementing the MH algorithm

We first need to build a function to calculate the mixture of normal distributions:

```
mixture_normal <- function(z, c, sigma_n){
  k <- length(c)
  zs <- matrix(NA, nrow = k, ncol = length(z))
  for (i in 1:k){
    zs[i,] <- rmvnorm(1, z, c[i] * sigma_n)
  }
  return(zs)
}
```